

Navigation

July 3, 2018

1 Navigation

In this notebook we train an agent in the Unity ML-Agents environment for the first project of the [Deep Reinforcement Learning Nanodegree](#). The agent code is mostly independent from the environment - we used a DQN agent developed to work with OpenAI Gym's LunarLander-v2 environment for the Deep Q-Network lesson. The agent and the corresponding networks are defined in a separate files.

1.0.1 1. Start the Environment

We begin by importing some necessary packages. If the code cell below returns an error, please revisit the project instructions to double-check that you have installed [Unity ML-Agents](#) and [NumPy](#).

```
In [1]: from unityagents import UnityEnvironment
        import numpy as np
```

Next, we will start the environment! *Before running the code cell below*, change the `file_name` parameter to match the location of the Unity environment that you downloaded.

- **Mac**: "path/to/Banana.app"
- **Windows** (x86): "path/to/Banana_Windows_x86/Banana.exe"
- **Windows** (x86_64): "path/to/Banana_Windows_x86_64/Banana.exe"
- **Linux** (x86): "path/to/Banana_Linux/Banana.x86"
- **Linux** (x86_64): "path/to/Banana_Linux/Banana.x86_64"
- **Linux** (x86, headless): "path/to/Banana_Linux_NoVis/Banana.x86"
- **Linux** (x86_64, headless): "path/to/Banana_Linux_NoVis/Banana.x86_64"

For instance, if you are using a Mac, then you downloaded `Banana.app`. If this file is in the same folder as the notebook, then the line below should appear as follows:

```
env = UnityEnvironment(file_name="Banana.app")
```

```
In [2]: env = UnityEnvironment(file_name="/data/Banana_Linux_NoVis/Banana.x86_64")
```

```

INFO:unityagents:
'Academy' started successfully!
Unity Academy name: Academy
    Number of Brains: 1
    Number of External Brains : 1
    Lesson number : 0
    Reset Parameters :

Unity brain name: BananaBrain
    Number of Visual Observations (per agent): 0
    Vector Observation space type: continuous
    Vector Observation space size (per agent): 37
    Number of stacked Vector Observation: 1
    Vector Action space type: discrete
    Vector Action space size (per agent): 4
    Vector Action descriptions: , , ,

```

Environments contain *brains* which are responsible for deciding the actions of their associated agents. Here we check for the first brain available, and set it as the default brain we will be controlling from Python.

```

In [3]: # get the default brain
        brain_name = env.brain_names[0]
        brain = env.brains[brain_name]

```

1.0.2 2. Examine the State and Action Spaces

The simulation contains a single agent that navigates a large environment. At each time step, it has four actions at its disposal: - 0 - walk forward - 1 - walk backward - 2 - turn left - 3 - turn right

The state space has 37 dimensions and contains the agent's velocity, along with ray-based perception of objects around agent's forward direction. A reward of +1 is provided for collecting a yellow banana, and a reward of -1 is provided for collecting a blue banana.

Run the code cell below to print some information about the environment.

```

In [4]: # reset the environment
        env_info = env.reset(train_mode=True)[brain_name]

        # number of agents in the environment
        print('Number of agents:', len(env_info.agents))

        # number of actions
        action_size = brain.vector_action_space_size
        print('Number of actions:', action_size)

        # examine the state space
        state = env_info.vector_observations[0]
        print('States look like:', state)

```

```
state_size = len(state)
print('States have length:', state_size)
```

Number of agents: 1

Number of actions: 4

```
States look like: [ 1.          0.          0.          0.          0.84408134  0.          0.
 1.          0.          0.0748472  0.          1.          0.          0.
 0.25755     1.          0.          0.          0.          0.74177343
 0.          1.          0.          0.          0.25854847  0.          0.
 1.          0.          0.09355672  0.          1.          0.          0.
 0.31969345  0.          0.          ]
```

States have length: 37

1.0.3 3. Start the Agent

```
In [5]: from agent import Agent
        agent = Agent(state_size=state_size, action_size=action_size, seed=0)
```

1.0.4 4. Train the Agent

When training the environment, set `train_mode=True`, so that the line for resetting the environment looks like the following:

```
env_info = env.reset(train_mode=True)[brain_name]
```

Importing packages Most of these are needed only for debugging and experimenting purposes. The agent is standalone.

```
In [6]: import torch
        from collections import namedtuple, deque

        #import numpy as np
        #import random
        #from model import QNetwork
        #import torch.nn.functional as F
        #import torch.optim as optim
```

Training parameters

```
In [7]: BUFFER_SIZE = int(1e5) # replay buffer size
        BATCH_SIZE = 64        # minibatch size
        GAMMA = 0.99           # discount factor
        TAU = 1e-3             # for soft update of target parameters
        LR = 5e-4               # learning rate
        UPDATE_EVERY = 4       # how often to update the network
```

Training algorithm and Agent training

```
In [8]: def dqn(n_episodes=2000, max_t=1000, eps_start=1.0, eps_end=0.01, eps_decay=0.995):
        """Deep Q-Learning.

        Params
        =====
        n_episodes (int): maximum number of training episodes
        max_t (int): maximum number of timesteps per episode
        eps_start (float): starting value of epsilon, for epsilon-greedy action selection
        eps_end (float): minimum value of epsilon
        eps_decay (float): multiplicative factor (per episode) for decreasing epsilon
        """

        scores = [] # list containing scores from each episode
        scores_window = deque(maxlen=100) # last 100 scores
        eps = eps_start # initialize epsilon
        for i_episode in range(1, n_episodes+1):
            env_info = env.reset(train_mode=True)[brain_name] # reset the environment
            state = env_info.vector_observations[0] # get the current state
            #state = env.reset()
            score = 0
            for t in range(max_t):
                action = agent.act(state, eps)
                env_info = env.step(action)[brain_name] # send the action to the env
                next_state = env_info.vector_observations[0] # get the next state
                reward = env_info.rewards[0] # get the reward
                done = env_info.local_done[0] # see if episode has finished
                agent.step(state, action, reward, next_state, done)
                state = next_state
                score += reward
                if done:
                    break
            scores_window.append(score) # save most recent score
            scores.append(score) # save most recent score
            eps = max(eps_end, eps_decay*eps) # decrease epsilon
            print('\rEpisode {} \tAverage Score: {:.2f}'.format(i_episode, np.mean(scores_window)))
            if i_episode % 100 == 0:
                print('\rEpisode {} \tAverage Score: {:.2f}'.format(i_episode, np.mean(scores_window)))
            if np.mean(scores_window) >= 14.0:
                print('\nEnvironment solved in {:d} episodes! \tAverage Score: {:.2f}'.format(i_episode, np.mean(scores_window)))
                torch.save(agent.qnetwork_local.state_dict(), 'checkpoint.pth')
                break
        return scores

#env_info = env.reset(train_mode=True)[brain_name]
scores = dqn()

Episode 100      Average Score: 1.00
Episode 200      Average Score: 4.52
```

```

Episode 300      Average Score: 8.26
Episode 400      Average Score: 9.80
Episode 500      Average Score: 11.98
Episode 600      Average Score: 13.07
Episode 627      Average Score: 14.01
Environment solved in 527 episodes!      Average Score: 14.01

```

```

In [9]: # save the trained weights to file
        torch.save(agent.qnetwork_local.state_dict(), 'checkpoint.pth')

```

Scores plot

```

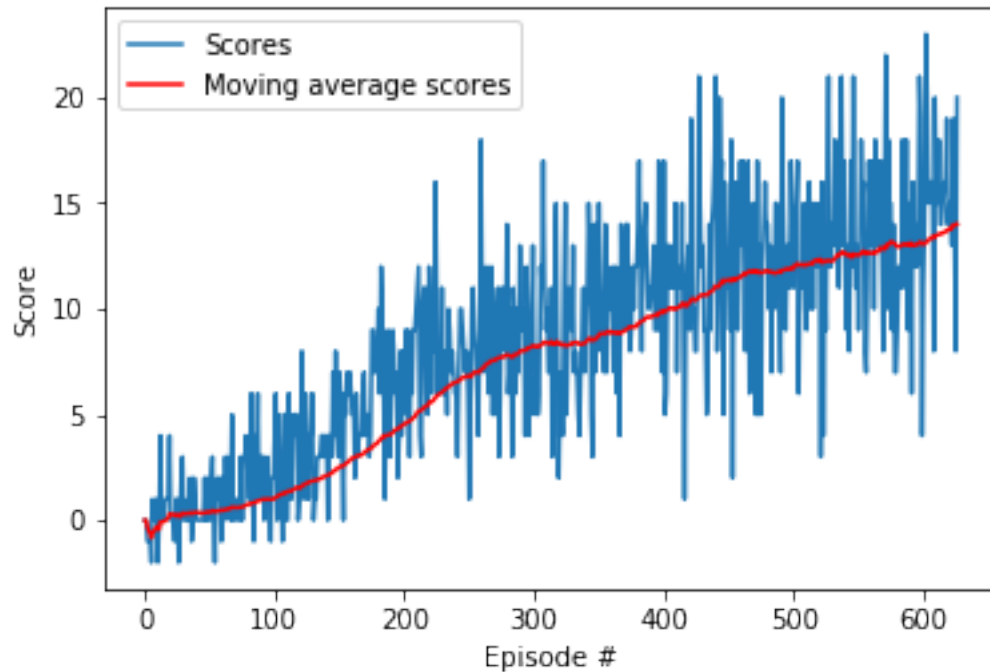
In [10]: import matplotlib.pyplot as plt
         %matplotlib inline

In [11]: # plot the scores
         fig = plt.figure()
         #ax = fig.add_subplot(111)
         plt.plot(np.arange(len(scores)), scores, label='Scores')
         plt.ylabel('Score')
         plt.xlabel('Episode #')

         cumsum_vec = np.cumsum(scores)
         window_width = 100
         ma_vec1 = np.cumsum(scores[:window_width])/(np.arange(window_width)+1)
         ma_vec2 = (cumsum_vec[window_width:] - cumsum_vec[:-window_width]) / window_width
         ma_vec = np.concatenate((ma_vec1, ma_vec2))
         plt.plot(np.arange(len(scores)), ma_vec, label='Moving average scores', color='r')
         plt.legend()
         #https://stackoverflow.com/a/34387987

         plt.show()

```



When finished, you can close the environment.

```
In [12]: env.close()
```

1.0.5 4. Watch a Smart Agent!

You can load the trained weights from file to watch a smart agent in the [Navigation_Demo.ipynb](#) !

1.0.6 References and Acknowledgments

This file based on the corresponding file in the first project of the Deep Reinforcement Learning Nanodegree, <https://www.udacity.com/course/deep-reinforcement-learning-nanodegree-nd893>

The agent and the training algorithms are from the *Deep Q-Networks* with slight modifications to account for the different environment (instead of OpenAI gym interface we use Unity environment).